

Avinash Reddy

Staff Site Reliability Engineer

Sunnyvale, CA avinashreddy.palem@gmail.com linkedin.com/in/avinashreddy30 areddy30.github.io

SUMMARY

Staff Site Reliability Engineer with 15 years of engineering experience, including 5+ years owning reliability for a multi-region Kubernetes platform across 100+ engineering teams. Deep experience defining SLIs, SLOs, and error budgets for critical request paths, leading incident response and blameless postmortems, and coaching engineering teams to own reliability. Hands-on expertise in Kubernetes, AWS, and Datadog observability tooling in high-throughput, low-latency environments. Built AI-powered incident automation that reduced operational toil by 40%+, applying generative AI to incident investigation, postmortem analysis, and runbook authoring. Drove cloud cost engineering that removed roughly \$1.5M from annual AWS spend. Background in multi-region, failure-isolation architecture. Certified Kubernetes Administrator (CKA).

EXPERIENCE

Staff Site Reliability Engineer

June 2021 – Present

Procore Technologies · Sunnyvale, CA

- Defined and monitored SLIs, SLOs, and error budgets based on EKS cluster health metrics and add-on component monitors across 30+ clusters serving 100+ engineering teams, making reliability measurable and tying error budgets to feature prioritization decisions.
- Led the incident lifecycle end to end — detection, response, communication, postmortem quality, and follow-up completion — for high-severity, customer-facing incidents across the production fleet, driving repeat incidents down through blameless postmortems.
- Specified repetitive Kubernetes operations as PRDs and encoded them as 20+ AI skills chaining Datadog, kubectl, and Terraform — rollout and upgrade testing, cluster lifecycle runbooks, incident investigation, and capacity/cost analysis — reducing operational toil by 40%+.
- Led reliability reviews for high-risk changes and new services (production readiness, capacity, failure modes, rollback), teaching through review rather than gatekeeping across 100+ engineering teams.
- Cut roughly \$1.5M in annual AWS spend: redesigned Karpenter node pool topology around Graviton, AMD, compute-optimized, GPU, and spot capacity; added ephemeral node pools so batch and transient workloads stopped holding long-lived capacity; and tiered cold EBS, EFS, and Sumo Logic data to infrequent-access storage (~\$1M compute, \$450K+ storage, six figures observability).
- Operated the core platform stack including Cilium (CNI), Istio (service mesh), and ArgoCD (GitOps), supporting continuous delivery and reliability for 100+ engineering teams.
- Designed and executed a SOX-compliant cluster decommissioning process covering governance, documentation, stakeholder coordination, and technical teardown, generalized into reusable frameworks for EKS clusters and standalone cloud assets (EC2, RDS, EBS).
- Mentored 3+ SRE engineers on Kubernetes operations, SRE principles, and on-call best practices — codifying production readiness checklists and on-call standards that persisted after stepping back.
- Led Terraform Infrastructure as Code rollouts across 30+ production clusters, enforcing cloud security standards and SOX-compliant governance.

Senior DevOps Engineer

January 2018 – June 2021

Intuit · Mountain View, CA

- Built a comprehensive observability stack using Datadog, Prometheus, and Grafana for monitoring, logging, and tracing across 1,000+ nodes, reducing mean time to detect incidents by 30%.
- Established structured root-cause analysis and blameless postmortem processes for high-severity incidents across voice, video, and chat production services, materially improving how the organization learned from incidents.
- Led Kubernetes adoption across the platform including OPA policy enforcement for cloud security, federated deployments, and custom resource definitions for distributed systems architecture.
- Architected CI/CD pipelines using Spinnaker, ArgoCD, Docker, and Kubernetes, automating deployments from commit to production.
- Migrated the Virtual Expert Platform to AWS within 60 days with zero downtime, implementing multi-region infrastructure for failure isolation.

DevOps / Site Reliability Engineer

2014 – 2018

Tata Consultancy Services · San Antonio, TX

- Automated fraud-detection infrastructure services using Chef, Ansible, Python, and Bash in an adversarial, real-time fraud detection domain, reducing deployment time 15x (from 3 weeks to 1 day).
- Improved infrastructure stability and reliability 35–70% by designing custom infrastructure services and leading blue/green deployment adoption with production readiness reviews.
- Maintained production services by measuring availability, latency, and system health, participating in on-call rotation and incident response.
- Administered identity governance and cloud security standards (RBAC, SSO via OKTA, SAML, OAuth) across Linux environments with TCP/IP networking.

Java/J2EE Developer & Production Support Engineer

2011 – 2014

Tata Consultancy Services · Chennai, India

- Automated fraud-detection batch pipelines using Java and Python across 15+ monthly jobs, reducing false positives 30%.
- Provided 24/7 L2 production support and on-call incident response for enterprise fraud, anti-money-laundering (AML), and financial-crime detection platforms, including case and alert management systems.
- Managed enterprise identity and access: provisioning and de-provisioning user and group access across target applications, with role-based access control and periodic access reviews for information-security compliance.
- Supported single sign-on and federated authentication using Okta, SAML, and OAuth, resolving access and federation incidents within SLA.

TECHNICAL SKILLS

Languages	Python, Bash, Java, SQL
Cloud Platforms	AWS (multi-region, multi-account), Azure, GCP
Cloud Cost / FinOps	Karpenter node pool design, Graviton and Spot adoption, storage tiering, log tiering
Kubernetes & Containers	EKS, Karpenter, Docker, Helm, Cilium, Istio, ArgoCD, Kargo
Infrastructure as Code	Terraform, Atlantis, CloudFormation, Chef, Ansible
Observability	Datadog, Prometheus, Grafana, Splunk, AWS CloudWatch
SRE Practices	SLIs, SLOs, Error Budgets, Incident Management, Blameless Postmortems, Production Readiness Reviews
AI & Reliability	AI-assisted Kubernetes operations, rollout testing, incident investigation, runbook authoring, generative-AI automation
Distributed Systems	Multi-region architecture, failure isolation, cascading failure prevention, load shedding
CI/CD	Jenkins, GitLab, Spinnaker, ArgoCD, GitHub Actions
Security & Networking	Linux, TCP/IP, OPA, Twistlock Defender, Cortex XDR, SOX Compliance, RBAC, SSO, SAML, OAuth
Data Stores	Redis, ElastiCache (failure modes and operations)

CERTIFICATIONS & EDUCATION

Certified Kubernetes Administrator (CKA)

Personal Leadership — Stanford University Graduate School of Business, 2019–2021

AWS Certified Solutions Architect

B.E. Electronics & Communications — Anna University

GitOps Certified for Argo — GitOps at Scale (Level II)